

Data Analytics

EEE 4774 & 6777

Module 4 - Classification

Deep Learning, Convolutional Neural Network

Spring 2022

Artificial Neural Network

- Origins in attempts to find mathematical representations of information processing in biological systems

(MLP) : fully connected neural network

- Also known as **Multilayer Perceptron** / Can be also regarded as **Multilayer Logistic Regression**

- Finds basis functions $\phi(x)$ adaptive to the training data


$$y(x) = f(c\phi(x)) = \sigma(w_2^T \phi(x))$$

$$\phi(x) = h(w_1^T x)$$

linear

nonlinear

raw data



Extensions

Mostly black-box approaches

Applications

Deep Neural Networks (Deep Learning), e.g.,

Supervised Learning

- Convolutional Neural Network (CNN) *image processing*
- Recurrent Neural Network (RNN) *Module 5 sequential data*
- Transformer (self-attention)
- Deep Generative Models *module 6*
- Deep Reinforcement Learning




Image Classification

- Object Recognition
- Action Recognition
- Face Recognition
- Speech Recognition
- Natural Language Processing *(NLP)*
- Video Understanding
- Time-series Prediction *Value Time*
- Anomaly Detection
- Robotics
- Autonomous Driving
- ...

Convolutional Neural Networks (CNN)

1D convolution



- Uses **convolution** in place of general matrix multiplication in at least one of the layers
- Convolution is a specialized kind of linear operation:

convolution operation


$$s(t) = x(t) * w(t) = \sum_{a=-\infty}^{\infty} x(a)w(t-a)$$

$$y[n] = x[n] * w[n] = \sum_{m=-\infty}^{\infty} w[m] x[n-m]$$

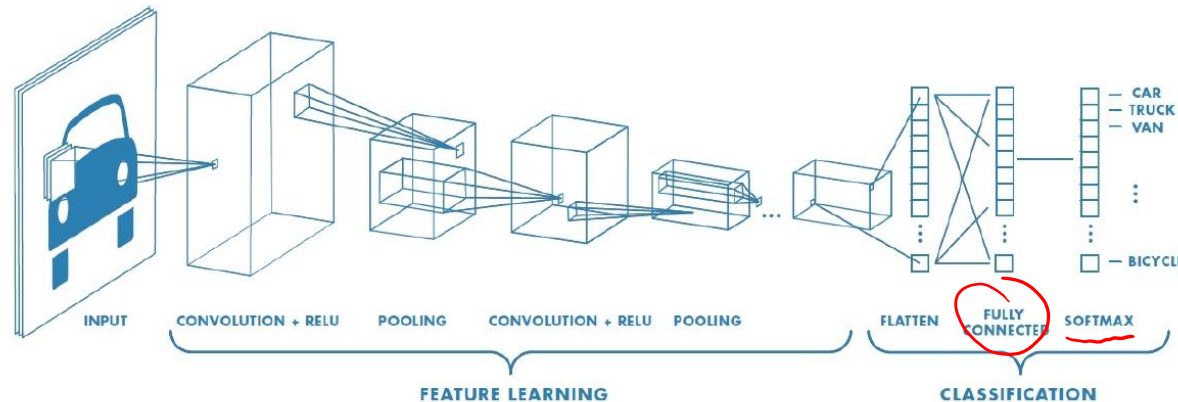
e.g.,

$$\hat{x}[n+1] = y[n] = w[0]x[n] + w[1]x[n-1] + w[2]x[n-2] + w[3]x[n-3] + w[4]x[n-4]$$


- Used for processing data that has a grid-like topology, e.g., time-series data (1-D grid), image data (2-D grid)
- Very successful in practice



$$W = \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_4 \end{bmatrix}$$



Convolution



| | | | | | | | | | |
|---|---|----|---|---|---|---|--|--|--|
| 0 | 0 | 0 | 0 | 0 | 0 | | | | |
| 0 | 5 | 10 | 2 | . | . | . | | | |
| 0 | 1 | 3 | 4 | . | . | . | | | |
| 0 | | | | | | | | | |
| 0 | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |


Padding

$$3 \times 5 + (-2)10 + (-1)3 = -8$$

16x16x3


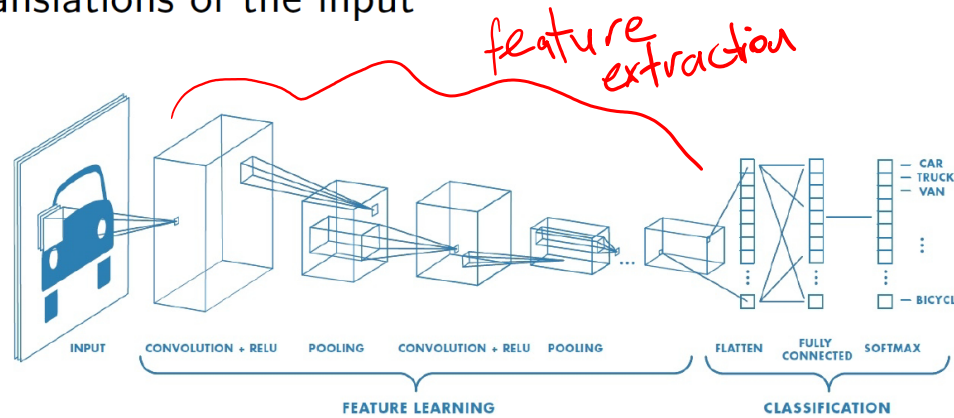
| | | | | | | | | | |
|---|---|---|---|---|---|--|--|--|--|
| 0 | 0 | 0 | 0 | 0 | 0 | | | | |
| 0 | . | . | . | . | . | | | | |
| 0 | . | . | . | . | . | | | | |
| 0 | | | | | | | | | |
| 0 | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |
| | | | | | | | | | |

Stride = 2



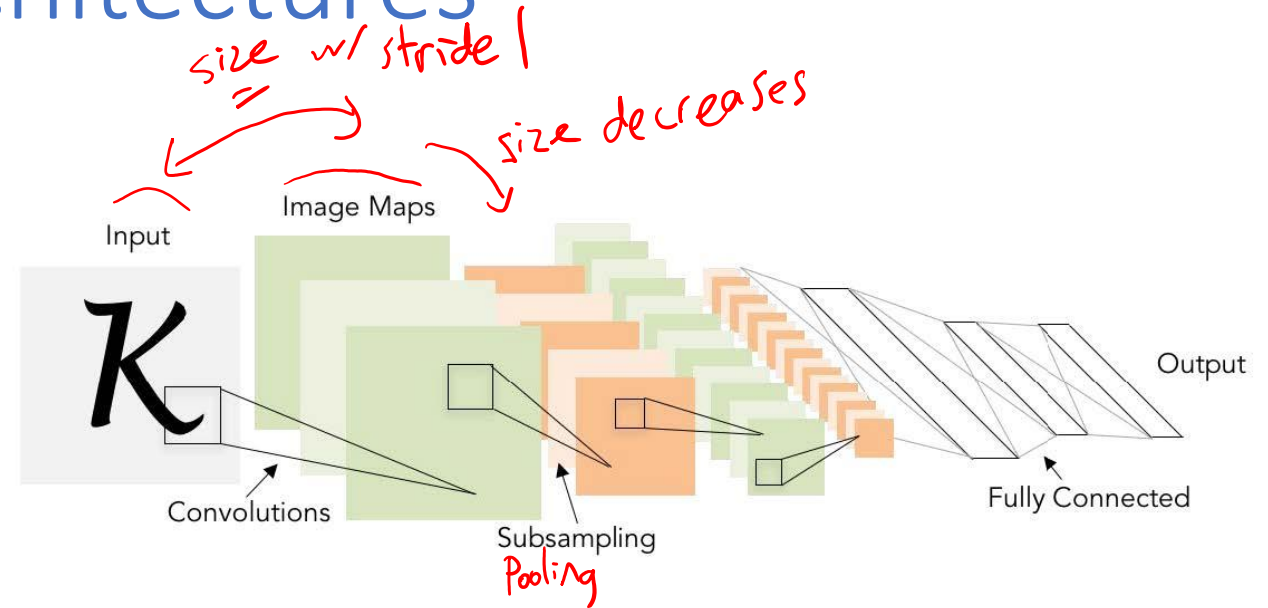
Convolutional Neural Networks

- 3 stages:
 - convolutional stage: linear activations
 - detector stage: nonlinear activation function such as Rectified Linear Unit (ReLU)
 - pooling stage: modify the output with a summary statistic of nearby outputs, e.g.,
- **Max Pooling:** reports the maximum output within a rectangular neighborhood
- **Average Pooling:** average of a rectangular neighborhood
- Pooling helps to make the representation approximately invariant (i.e., robust) to small translations of the input



CNN Architectures

- LeNet-5 [LeCunn et al. 1998]
- AlexNet [Krizhevsky et al. 2012]
- VGG [Simonyan et al. 2014]
- GoogLeNet [Szegedy et al. 2015]
- ResNet [He et al. 2016]
- ...
- ...




LeNet-5: Conv filters were 5x5, applied at stride 1
Subsampling (Pooling) layers were 2x2 applied at stride 2
i.e. architecture is [CONV-POOL-CONV-POOL-FC-FC]

4 layers

Rise of Deep Learning

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



First Percentage

GPU


AlexNet

Img size = $224 \times 224 \times 3$
GPU 2 (48 filters/neurons)

Architecture:

CONV1
MAX POOL1
NORM1
CONV2
MAX POOL2
NORM2
CONV3
CONV4
CONV5
Max POOL3


FC6
FC7
FC8




Filter size = 11x11
GPU 1
filters
neurons

2 GPU

Max Pooling



AlexNet



Trained on 2 GPUs.
Half the neurons on
each GPU.

Connections only
within the same
GPU

Connections
across GPUs

AlexNet

- First use of ReLU
- Used Norm layers
- **Data augmentation (overfitting)**
- **Dropout (overfitting)**
 - Randomly drop neurons for each training instance in feedforward and backpropagation with probability 0.5
 - “Every time an input is presented, the neural network samples a different architecture, but all these architectures share weights”
 - “At test time, use all the neurons but multiply their outputs by 0.5”
- SGD for weight update in training
 - Gradient batch size 128
 - Momentum 0.9
 - Learning rate 1e-2, reduced by 10 manually when validation accuracy plateaus
- 7 CNN ensemble: 18.2% -> 15.4%

Nonlin. activation func.
increase diversity in training dataset

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca


Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca

Abstract

We trained a large, deep convolutional neural network to classify the 1.2 million high-resolution images in the ImageNet LSVRC-2010 contest into the 1000 different classes. On the test data, we achieved top-1 and top-5 error rates of 37.5% and 17.0% which is considerably better than the previous state-of-the-art. The neural network, which has 60 million parameters and 650,000 neurons, consists of five convolutional layers, some of which are followed by max-pooling layers, and three fully-connected layers with a final 1000-way softmax. To make training faster, we used non-saturating neurons and a very efficient GPU implementation of the convolution operation. To reduce overfitting in the fully-connected layers we employed a recently-developed regularization method called “dropout” that proved to be very effective. We also entered a variant of this model in the ILSVRC-2012 competition and achieved a winning top-5 test error rate of 15.3%, compared to 26.2% achieved by the second-best entry.

ResNet



ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



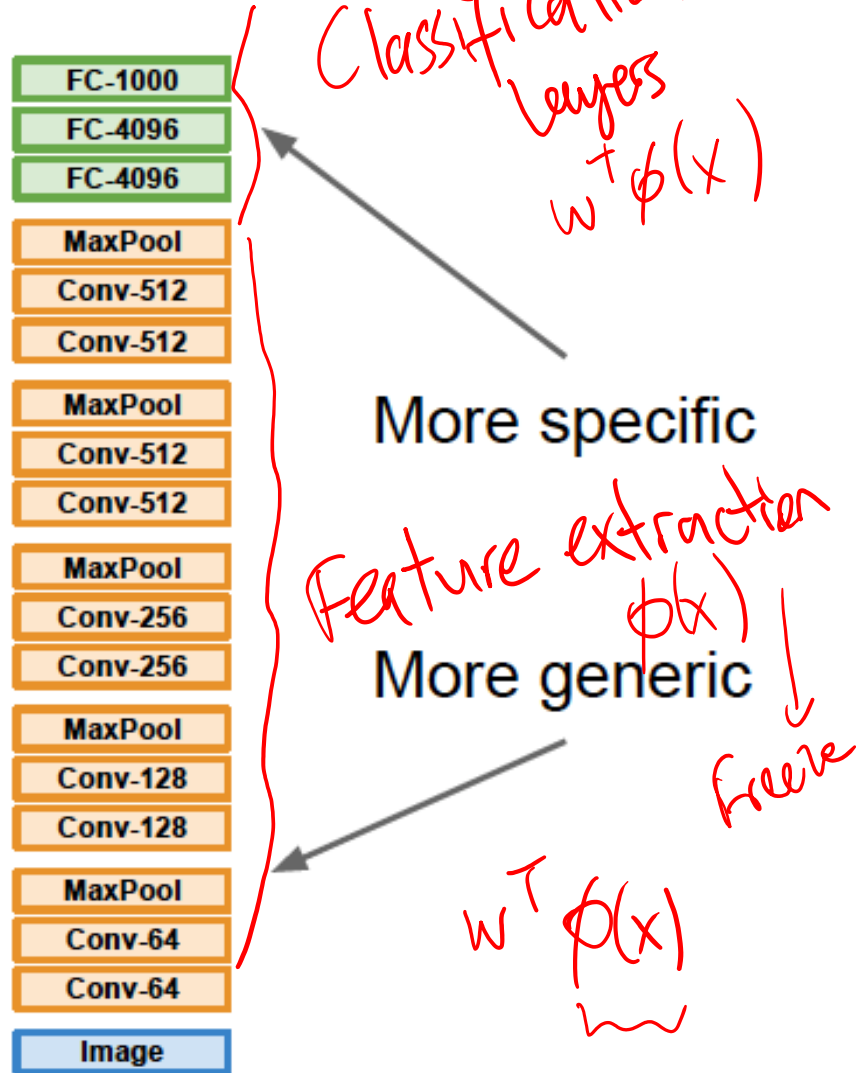
ResNet

Residual

- Very deep networks using residual connections
- **152-layer** model for ImageNet
 - Outperformed the human-level performance
- Now focus shifted to Efficient Networks:
 - Lots of tiny networks aimed at mobile devices: MobileNet, ShuffleNet, etc.



Transfer Learning



| | | |
|----------------------------|---|--|
| | very similar dataset | very different dataset |
| very little data | Finetune linear classifier on top layer | <u>You're in trouble...</u> Try data augmentation / collect more data |
| quite a lot of data | Finetune a <u>few layers</u> | Finetune a <u>larger number of layers</u> |

Transfer Learning with CNNs

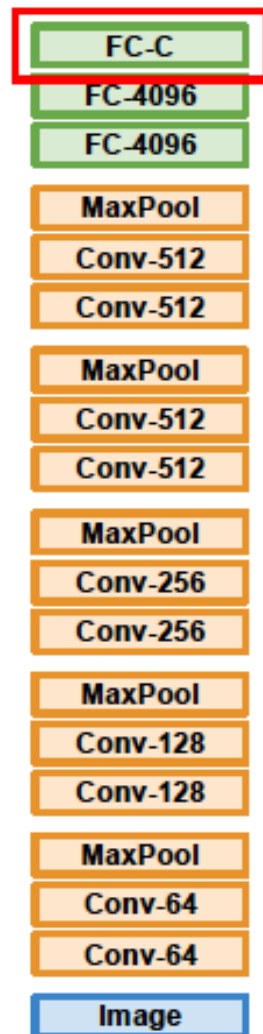
Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014
Razavian et al, "CNN Features Off-the-Shelf: An Astounding Baseline for Recognition", CVPR Workshops 2014

1. Train on Imagenet



↓
Pre-trained
model

2. Small Dataset (C classes)



Reinitialize
this and train

Freeze these

3. Bigger dataset



Train these

With bigger
dataset, train
more layers

Freeze these

Lower learning rate
when finetuning;
1/10 of original LR
is good starting
point

↓
0.01

Transfer Learning

In practice:

- Take a pretrained model
 - Trained on a very large dataset such as ImageNet
 - “Model Garden” of pretrained models:
<https://github.com/tensorflow/models>
<https://github.com/pytorch/vision>
- Train only a few last layers on your dataset