

Data Analytics

EEE 4774 & 6777

Module 4 - Classification

Deep Learning, Convolutional Neural Network


Spring 2022

Artificial Neural Network

- Origins in attempts to find mathematical representations of information processing in biological systems
- Also known as *Multilayer Perceptron* / Can be also regarded as *Multilayer Logistic Regression*
- Finds basis functions $\phi(\mathbf{x})$ adaptive to the training data

$$y(\mathbf{x}) = f(c\phi(\mathbf{x})) = \sigma(\mathbf{w}_2^T \phi(\mathbf{x}))$$

$$\phi(\mathbf{x}) = h(\mathbf{w}_1^T \mathbf{x})$$



Extensions

Deep Neural Networks (Deep Learning), e.g.,

- Convolutional Neural Network (CNN)
- Recurrent Neural Network (RNN)
- Transformer (self-attention)
- Deep Generative Models
- Deep Reinforcement Learning

Applications



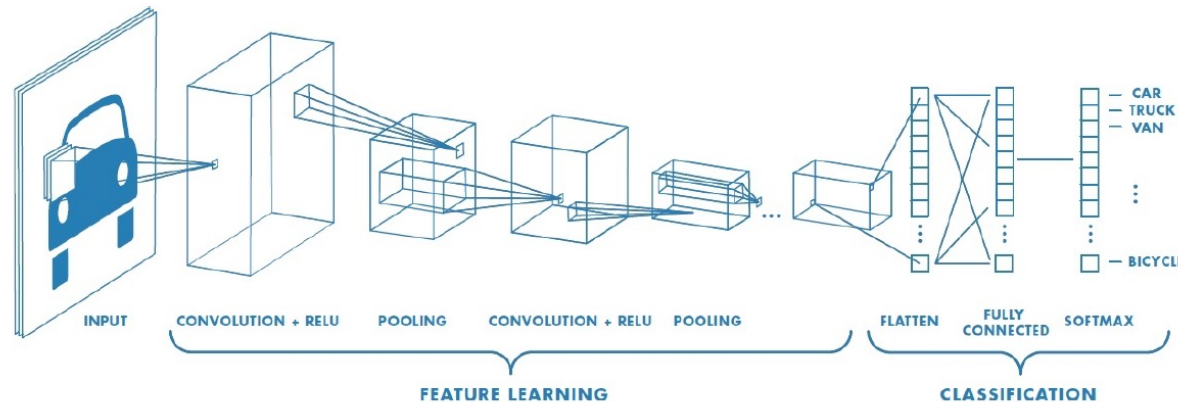
- Object Recognition
- Action Recognition
- Face Recognition
- Speech Recognition
- Natural Language Processing
- Video Understanding
- Time-series Prediction
- Anomaly Detection
- Robotics
- Autonomous Driving
- ...

Convolutional Neural Networks


- Uses **convolution** in place of general matrix multiplication in at least one of the layers
- Convolution is a specialized kind of linear operation:

$$s(t) = x(t) * w(t) = \sum_{a=-\infty}^{\infty} x(a)w(t-a)$$

- Used for processing data that has a grid-like topology, e.g., time-series data (1-D grid), image data (2-D grid)
- Very successful in practice



Convolution



0	0	0	0	0	0			
0								
0								
0								
0								

Padding

0	0	0	0	0	0			
0								
0								
0								
0								
0								

Stride

Convolutional Neural Networks

- 3 stages:
 - convolutional stage: linear activations
 - detector stage: nonlinear activation function such as Rectified Linear Unit (ReLU)
 - pooling stage: modify the output with a summary statistic of nearby outputs, e.g.,
- **Max Pooling:** reports the maximum output within a rectangular neighborhood
- **Average Pooling:** average of a rectangular neighborhood
- Pooling helps to make the representation approximately invariant (i.e., robust) to small translations of the input

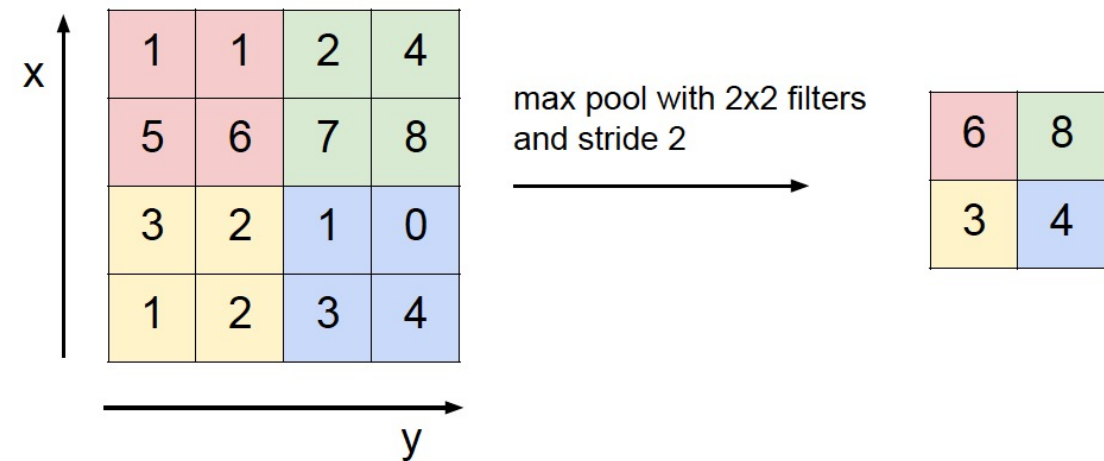
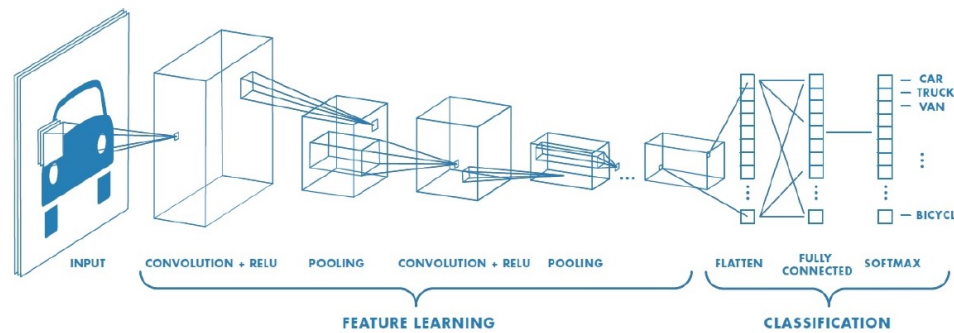



Image courtesy of Fei-Fei Li, Ranjay Krishna, Danfei Xu

CNN Architectures


- LeNet-5 [LeCunn et al. 1998]
- AlexNet [Krizhevsky et al. 2012]
- VGG [Simonyan et al. 2014]
- GoogLeNet [Szegedy et al. 2015]
- ResNet [He et al. 2016]
- ...
- ...



LeNet-5: Conv filters were 5x5, applied at stride 1
Subsampling (Pooling) layers were 2x2 applied at stride 2
i.e. architecture is [CONV-POOL-CONV-POOL-FC-FC]

Rise of Deep Learning

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



AlexNet

Architecture:

CONV1

MAX POOL1

NORM1

CONV2

MAX POOL2

NORM2

CONV3

CONV4


CONV5

Max POOL3


FC6

FC7

FC8



AlexNet



Trained on 2 GPUs.
Half the neurons on
each GPU.

Connections only
within the same
GPU

Connections
across GPUs

AlexNet

- First use of ReLU
- Used Norm layers
- **Data augmentation (overfitting)**
- **Dropout (overfitting)**
 - Randomly drop neurons for each training instance in feedforward and backpropagation with probability 0.5
 - “Every time an input is presented, the neural network samples a different architecture, but all these architectures share weights”
 - “At test time, use all the neurons but multiply their outputs by 0.5”
- SGD for weight update in training
 - Gradient batch size 128
 - Momentum 0.9
 - Learning rate 1e-2, reduced by 10 manually when validation accuracy plateaus
- 7 CNN ensemble: 18.2% -> 15.4%

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca


Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca

Abstract

We trained a large, deep convolutional neural network to classify the 1.2 million high-resolution images in the ImageNet LSVRC-2010 contest into the 1000 different classes. On the test data, we achieved top-1 and top-5 error rates of 37.5% and 17.0% which is considerably better than the previous state-of-the-art. The neural network, which has 60 million parameters and 650,000 neurons, consists of five convolutional layers, some of which are followed by max-pooling layers, and three fully-connected layers with a final 1000-way softmax. To make training faster, we used non-saturating neurons and a very efficient GPU implementation of the convolution operation. To reduce overfitting in the fully-connected layers we employed a recently-developed regularization method called “dropout” that proved to be very effective. We also entered a variant of this model in the ILSVRC-2012 competition and achieved a winning top-5 test error rate of 15.3%, compared to 26.2% achieved by the second-best entry.



ResNet

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



ResNet

- Very deep networks using residual connections
- **152-layer** model for ImageNet
 - Outperformed the human-level performance
- Now focus shifted to Efficient Networks:
 - Lots of tiny networks aimed at mobile devices: MobileNet, ShuffleNet, etc.



Transfer Learning



More specific


More generic

	very similar dataset	very different dataset
very little data	Finetune linear classifier on top layer	You're in trouble... Try data augmentation / collect more data
quite a lot of data	Finetune a few layers	Finetune a larger number of layers

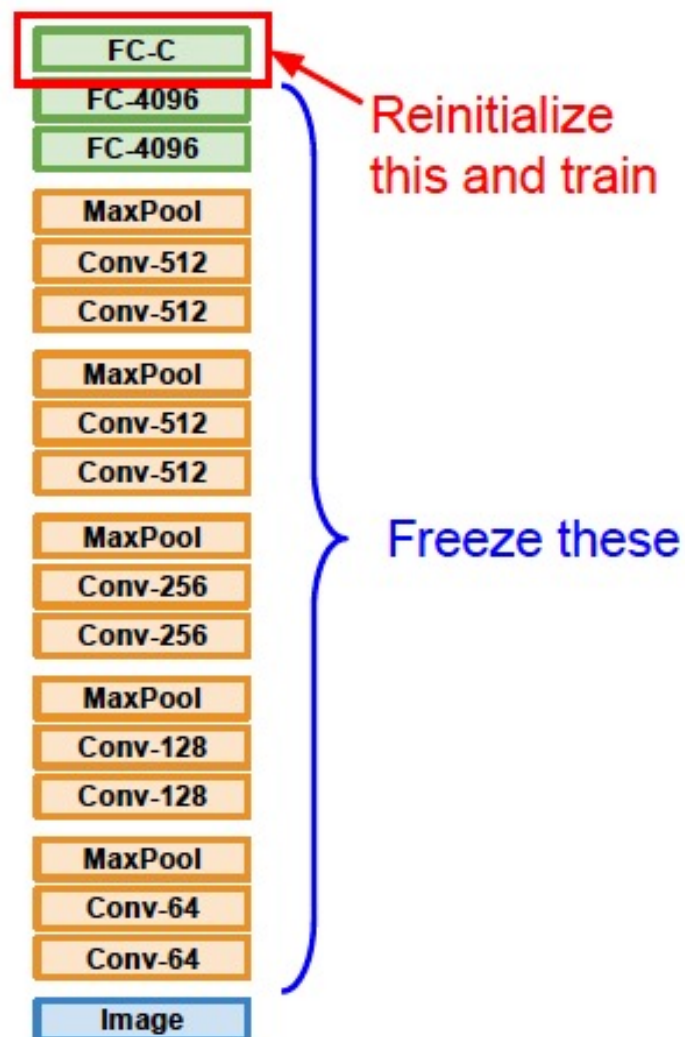
Transfer Learning with CNNs

Donahue et al, "DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition", ICML 2014
Razavian et al, "CNN Features Off-the-Shelf: An Astounding Baseline for Recognition", CVPR Workshops 2014

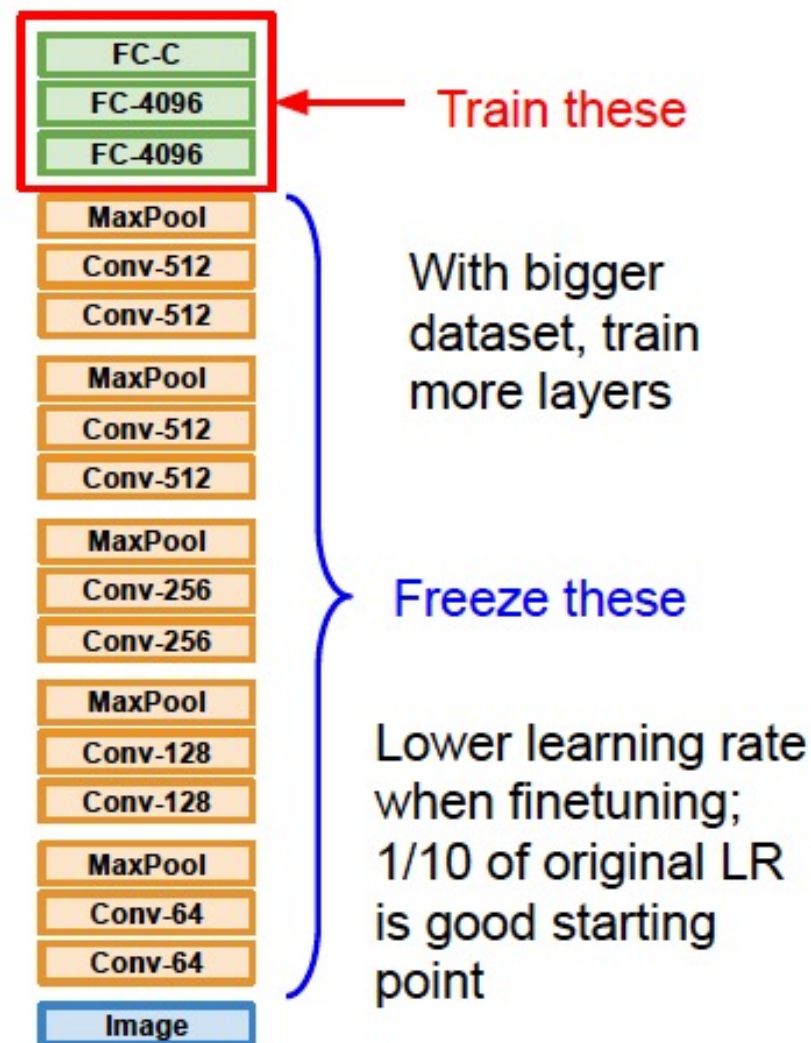
1. Train on Imagenet



2. Small Dataset (C classes)



3. Bigger dataset



Transfer Learning

In practice:

- Take a pretrained model
 - Trained on a very large dataset such as ImageNet
 - “Model Garden” of pretrained models:
<https://github.com/tensorflow/models>
<https://github.com/pytorch/vision>
- Train only a few last layers on your dataset