


Data Analytics Final Project: Image Classification

Ahmed Shahabaz
U 8941-5490


Problem Statement

- Classify test images correctly
- Steps:
 - Input pre-processing
 - Feature extraction - (CNN)
 - Classification
- Challenges:
 - Imbalanced dataset
 - Different lighting conditions in images

Dataset (train) distribution



Dataset (train) distribution using [*`torch.utils.data.WeightedRandomSampler`*](#)



Dataset Statistics

```
Data loader in train mode!  
Number of classes: 19.  
Train Data size: 2823.  
Validation Data size: 647.  
-----  
  
Data loader in test mode!  
Number of classes: 19.  
Test Data size: 570.  
-----
```

Problems of using `torch.utils.data.WeightedRandomSampler`

- Unbalanced data issue is solved using oversampling
 - Images are repeated for classes with lower number of images
 - So model might overfit to those classes. As model sees same sample multiple times instead of only once in an epoch
 - Solution: Use Data Augmentation/Image Transformation
- Another approach can be weighted loss

Different data augmentations


```
if mode == "train" or mode == "Train":
    transforms_list = [
        transforms.Resize((args.image_size)),
        transforms.CenterCrop(args.crop_size),
        transforms.RandomHorizontalFlip(),
        transforms.RandomVerticalFlip(),
        #transforms.RandomAutocontrast(),
        #transforms.RandomEqualize(),
        transforms.RandomApply(torch.nn.ModuleList([
            transforms.ColorJitter(),
            transforms.GaussianBlur(15),
        ]), p = 0.3),
        transforms.RandomRotation(20),
        transforms.ToTensor(),
        transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225))
    ]

else:
    transforms_list = [
        transforms.Resize((args.image_size)),
        transforms.CenterCrop(args.crop_size),
        transforms.ToTensor(),
        transforms.Normalize((0.485, 0.456, 0.406), (0.229, 0.224, 0.225))
    ]
```


Addressing overfitting

- Early Stopping
 - If no change in validation loss for a few epochs then training stops.
- Different regularization
 - Weight decay
 - Multiply the sum of squares of weights (parameter) with another smaller number. This number is called **weight decay** or wd.
 - Prevents weights getting too large
 - $wd = 1e-6$ worked better for our case
 - Drop out
 - Randomly drops some of the neurons of a layer
 - Only tried 0.5


ResNet18: Loss w/o pre-training (no lr. scheduler)




ResNet18: Accuracy w/o pre-training (no lr. scheduler)




ResNet34: Accuracy w/o pre-training (lr. scheduler = 50)




ResNet34: Accuracy with pre-training (no lr. scheduler)



ResNet34 (fc layer): Accuracy with pre-training (lr. scheduler = 15)




ResNet18: Accuracy with pre-training and (lr. scheduler = 10)



Comments on results of Deep Learning

- ResNet34, ResNet50, ResNet101 and VGG16 with pre-trained weights are better
 - Only the fc layer of the ResNets and the classification layer of VGG16 was trained
 - Using lr. scheduler = 15 gives better performance
- Overfits easily to ResNet18
- Optimizer used: SGD

VGG16 (classification layer): Loss using ADAM optimizer (lr. scheduler = 15)



Test Results

```
Loading saved model weights...

*** Creating predictions for each model for ensemble ***

Model: ResNet18
Epoch 0,Test: 100%|████████████████████████████████████████| 18/18 [00:01<00:00, 9.70batch/s, Acc=0.93, Lss=0.22]

Model: VGG16_SGD
Epoch 0,Test: 100%|████████████████████████████████████████| 18/18 [00:03<00:00, 5.22batch/s, Acc=0.94, Lss=0.23]

Model: VGG16_adam
Epoch 0,Test: 100%|████████████████████████████████████████| 18/18 [00:03<00:00, 5.28batch/s, Acc=0.93, Lss=0.8]

Model: ResNet34
Epoch 0,Test: 100%|████████████████████████████████████████| 18/18 [00:01<00:00, 9.75batch/s, Acc=0.95, Lss=0.18]

Model: ResNet50
Epoch 0,Test: 100%|████████████████████████████████████████| 18/18 [00:02<00:00, 7.41batch/s, Acc=0.96, Lss=0.14]


Model: ResNet50
Epoch 0,Test: 100%|████████████████████████████████████████| 18/18 [00:02<00:00, 7.31batch/s, Acc=0.96, Lss=0.14]

Model: ResNet101
Epoch 0,Test: 100%|████████████████████████████████████████| 18/18 [00:03<00:00, 5.35batch/s, Acc=0.96, Lss=0.13]


*** Performance of the Ensemble (Majority Voting) of 6 models ***

Final Test accuracy after ensemble is: 97.71929824559689
```

Confusion matrix of ensemble vs ResNet101 (fc layer)



CM of ensemble of deep learners



CM of ResNet101

RandomForest

- RandomForestClassifier from sklearn.ensemble

n-estimators	50	100	200	300	500
Accuracy	0.17543	0.21754	0.20350	0.22280	0.22807

Support Vector Machine (SVM)

- sklearn implementation of SVM

SVM Kernel	Linear	RBF
Accuracy	0.13333	0.18070

Conclusion

- Deep learning algorithms outperformed traditional algorithms
- Ensemble of multiple deep learners was better than any single one of them
- In case of deep learning transfer learning worked much better than training from scratch
 - Which says our training dataset was not large enough for training a deep learning algorithm from scratch

Thank You